

CS 291 Guest Lecture

Andrew Mutz
Chief Technology Officer, Property Management
November 14, 2017



Outline for today

Appfolio Overview

Technology Overview

Engineering Overview

Scalability Lessons Learned

Challenges Ahead





The story starts with the company and the product. (Why?)

Appfolio makes

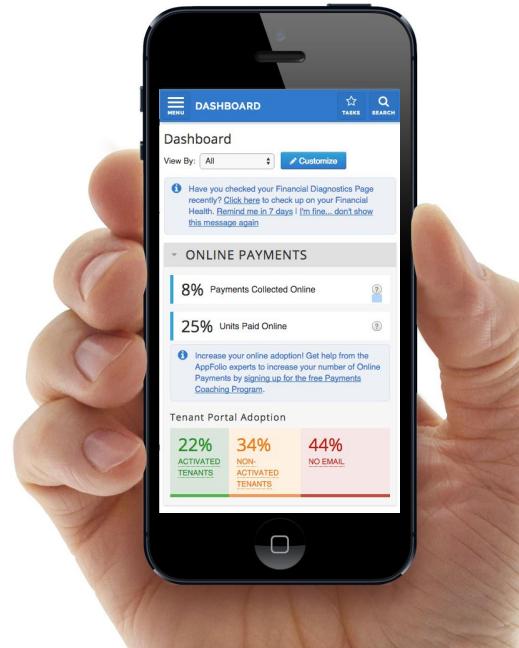
- Business software, sold as a service
- For small and medium sized businesses
- In specific industries



“Business software, sold as a service”

- Software is referred to as “a service” (SaaS) when the software is managed and run by an outside company
 - Think Outlook vs. GMail
 - Business software used to be primarily on the premises of the customer (and still is to a great deal)
 - Customer data is stored in data centers, managed by experts
 - Users interact with the system through web browsers and mobile apps
 - Updates are frequent
 - Users pay per month
- Software as a Service used to be the exception, but it is increasingly the norm
 - Consumer adoption of SaaS is very high: GMail, Facebook, Snapchat, YouTube, etc
 - Business adoption also high, but transition is slower

The screenshot shows the AppFolio Property Management dashboard. At the top, there's a navigation bar with links for DASHBOARD, LEASING, PROPERTIES, PEOPLE, ACCOUNTING, MAINTENANCE, REPORTING, and a search bar. A message at the top left encourages checking financial diagnostics. Below the navigation is a section titled 'ONLINE PAYMENTS' showing 8% Payments Collected Online and 25% Units Paid Online. Another section, 'Tenant Portal Adoption', shows 22% ACTIVATED TENANTS, 34% NON-ACTIVATED TENANTS, and 44% NO EMAIL. On the right side, there's a sidebar with links for PROPERTY (New Property), PEOPLE (Move In Tenant, Tenant Receipt, Enter Bill, New Owner, New Vendor, Rent Increase, Enter Utility Bill), and REPORTS (Delinquency, Tenant Ledger, Income Statement, Unit Vacancy Detail, Rent Roll).





"For small and medium sized businesses"

- Employees: Tens or hundreds
- Appfolio Property Manager doesn't serve either end of the spectrum:
 - If someone manages less than 50 units, they might not be a good fit
 - If someone manages many tens of thousands of units, they might not be a good fit
- SMB is not just about the number of users, it is an approach to building software
 - Customers are large enough that they are experts in their business
 - Customers are small enough that they do not need their software customized
 - Large companies tend to prefer software that is customized to their particular business practices
 - Appfolio does not build features or customize the product for a single company.
 - Caveat: we can turn features on and off for customers

“In specific industries”

- Software that is targeted at a specific industry is referred to as “Vertical” software
- Instead of picking a specific type of work that is universal to all companies, we build exactly what is needed by a specific type of company
- Vertical software emphasizes integration between different functional areas
 - The CRM system is integrated with the Accounting system, Calendaring functionality is integrated with the leasing system, etc
- Building vertical products generally provides a UX that is easier to use and of higher value to the user.
- Appfolio’s mission is to provide a portfolio of applications to a variety of verticals
 - Today we serve Property Management and Legal
 - Tomorrow many more





Appfolio Company Highlights

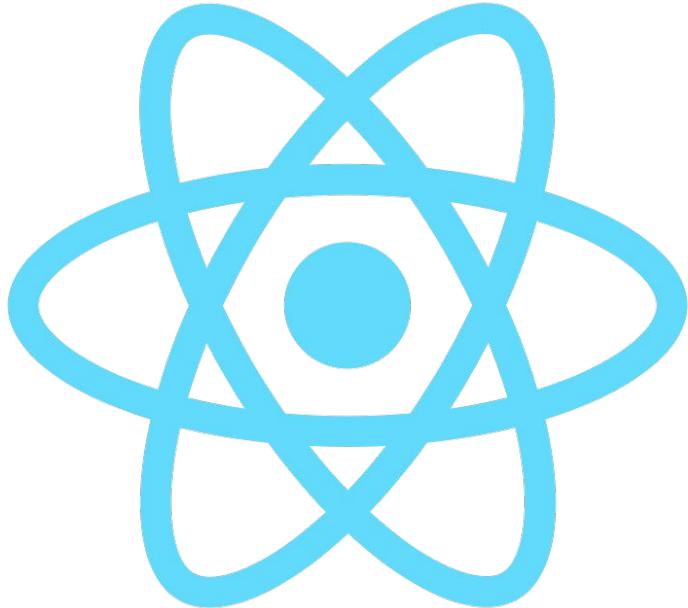
- Founded in 2006 by Klaus Schauser and Jon Walker
- About 600 employees, of those about 100 are engineers
- Headquartered in Santa Barbara with offices in San Diego and Dallas
- Number of customers
 - Property Management: ~10,000 companies
 - Legal: ~10,000 companies
- Number of “units” managed by Appfolio: ~3 Million
- Electronically transact ~\$20 billion dollars in rent annually
- Facilitate background check on ~2 million people annually
- ~20 million HTTP requests a day
- Publicly traded, market cap ~\$1.5 billion



Technology Overview

- Front end
- Application Layer
- Data Layer
- Operating System
- Hardware





Technology Overview: Front End

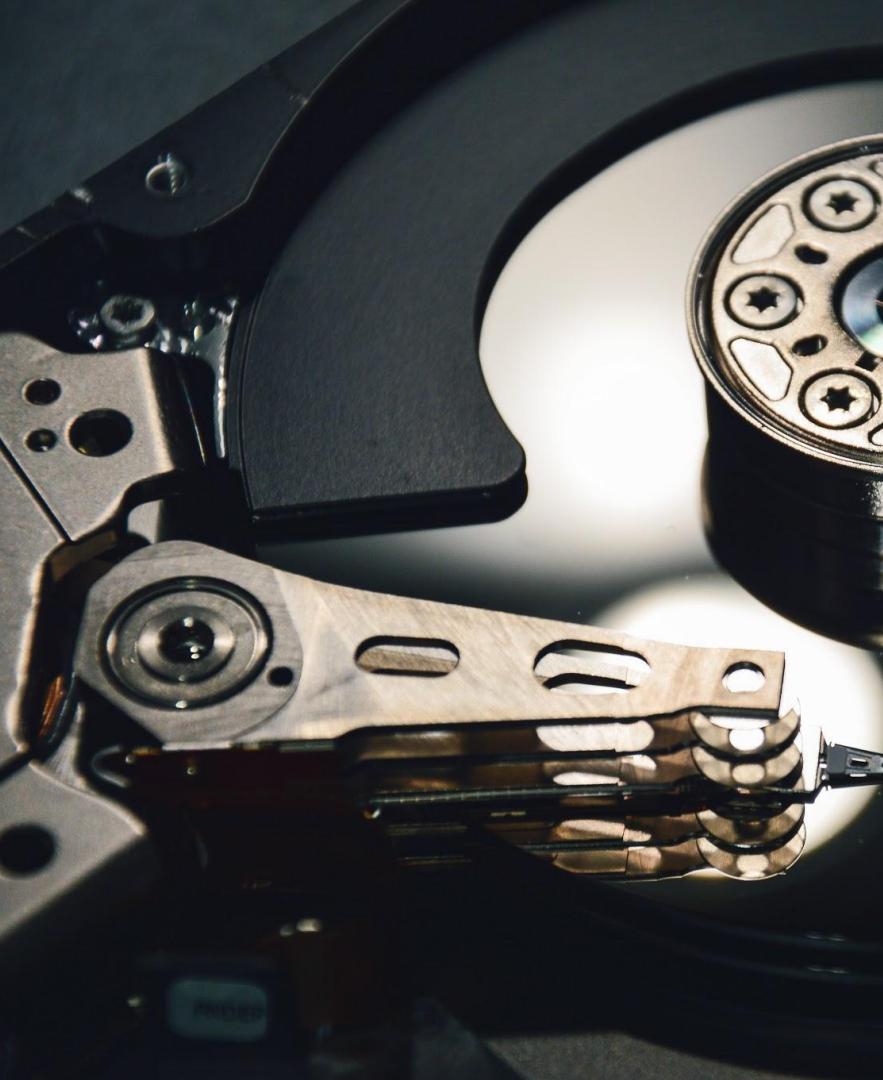
- Traditionally markup is generated on the server, but increasingly on the client
 - Client-side UIs built in React
 - We use MobX and Redux to manage React state
 - REST/JSON to communicate to back end
 - ES2015, ECMA stage 2
- Mobile apps built with React Native
 - Allows devs to write native apps using javascript and the same framework as client-side web
- Mobile apps use web views and responsive web design for many parts of the UI
 - This allows our developers to write a feature once and have a high quality experience across traditional web, mobile web, iOS app and Android app
 - Web views aren't perfect, but this solution is perfect for our product and business
- Styling is done using a bootstrap-style shared component library



Technology Overview: Application Layer

- All of our applications are built using Ruby on Rails
- Why Ruby?
 - Languages can be thought of as on a spectrum, productivity vs. performance
 - The right place to be on this spectrum is determined by your product and company
 - Appfolio users are high revenue-per-request, so the cost of operating servers is a tiny fraction of the revenue they generate
- Why Rails?
 - High developer productivity: convention over configuration allows us to go fast
- App servers run phusion passenger
 - Passenger supports multi-threaded, but we do multi-process. Why?
- Nginx sits in front of passenger, but passenger speaks HTTP. Why do we need both?





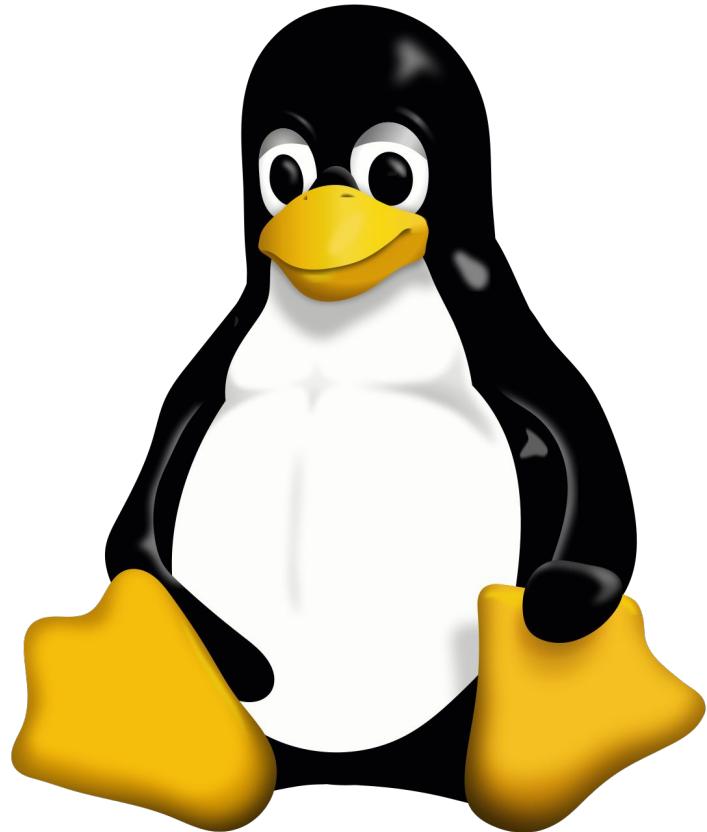
Technology Overview: Data Layer

- MySQL for relational data
 - Each customer's data is stored in a separate relational database
- Memcached for caching
 - Caching used extensively to reduce trips to the database
 - Permissions, frequently used accounting data, etc
- RabbitMQ as a message bus
 - Our asynchronous messaging is primarily used for asynchronous jobs
 - We don't generally rely on asynchronous messaging between services
- S3 for large file storage
 - No customer limit on storage
 - Users can attach arbitrary files to entities
 - As system of record, this gets used extensively
- SOLR for search
 - Application uses text search extensively (global search, entity selection)
 - SOLR uses Lucene under the hood
 - Creative use of ORM & SQL to get high performance SOLR-MySQL sync



Technology Overview: Operating System

- All our servers run CentOS Linux
- Transitioning to run our processes inside containers.
 - What are containers?
 - Filesystem and network isolation for unix processes
 - Why containers?
 - Increased isolation between apps
 - Developers can build in any language
 - Increased reproducibility between prod and dev
 - Docker vs. Rocket
 - Orchestration
 - Kubernetes vs. docker compose vs. nothing





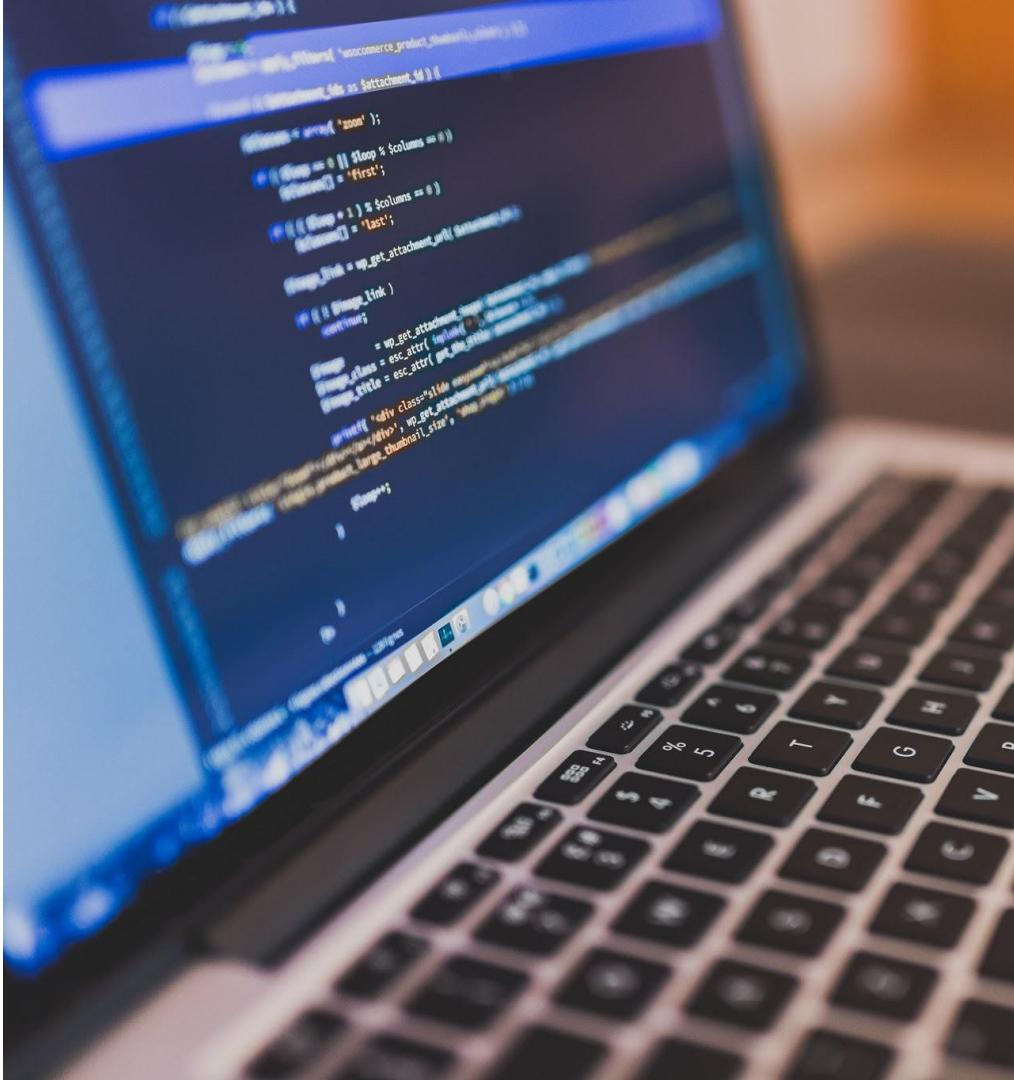
Technology Overview: Hardware

- APM runs in server colocation facilities
 - Dallas and Denver
 - ~200 servers
 - For load balancing and request routing we use Citrix Netscaler
 - For quick (but not cheap) durability, we use NetApp Filer
 - For load balancers, routers, switches, and filer we have two of everything
 - If any one of these devices fails, we can run on the backup until we can physically replace the failed hardware
 - For application servers and databases, we have two possible failover solutions
 - We can bring up the app servers and databases on another machine in the local datacenter
 - Or we can fail over to the twin machine in the remote data center
 - If a datacenter fails, we failover all machines to the other datacenter
 - How do clients know to route requests to Denver after a failover instead of to Dallas?
- Dell PowerEdge R630
 - Intel Xeon E5-2650 2.2GHz, Hyper-threading, 12 Cores, 24 Threads
 - RAM: 128GB
 - Disks: 3 hot-swappable 1.2GB, 10k rpm
- MyCase is in AWS.
 - Why is it different?



Engineering Overview

- TDD
- Agile/Scrum
- Release Cycle





Engineering Overview: Test-driven Design

- Test-driven development:
 - Write the test (watch it fail)
 - Write the code (watch it pass)
 - Refactor (tests should pass)
- Goal: by only writing code that has a corresponding test, you are guaranteed good test coverage.
- Third step (refactor) is very important and frequently forgotten
 - Without it, you get a very organic design.
- AppFolio has an extremely large rails app (~300K lines of code)
 - Heavily tested code is extremely valuable (we have more test code than actual code)
 - Currently releasing on a cadence that would be impossible without large test harness



Engineering Overview: Test-driven design

- TDD advocates can be dogmatic about test-first, we are more pragmatic
 - Since the goal is high test coverage, the order we do this is not important as long as we get high coverage
- AppFolio tests at every level:
 - Unit tests: Is a given object or model class working as designed?
 - Functional Tests: Is a web request working as designed?
 - Selenium Tests: Is a sequence of browser actions across the app working as designed?
 - Page objects keep this manageable
 - JavaScript Tests: Are our Javascript methods behaving correctly
 - RPC integration tests: Are our separately deployed services interacting correctly?
 - Saved report changes: Is this version of the application compatible with previously saved data?





Engineering Overview: TDD

- Integrating disparate changes is hard if you do it rarely
- Basic idea: check in often and run entire test suite each time
 - Goal: developers instantly know when they are introducing bugs into the system and can fix immediately
- All features are built on feature branches. CI evaluates these branches before merging.
 - Developers merge their work to master as frequently as possible.
- Continuous Integration is the backbone of our development process
 - Spend a lot of time and money to make it fast and easy
 - Automated tests are so important, we have more test code than production code
 - Automated tests are so important, we have about as many CI servers as production servers
- Each time a developer checks in, tests are run
 - Automatically deployed
 - Testing servers automatically migrated



Engineering Overview: Release cycle

- We release daily to a small group of customers
- We release every week to all customers
- Technically, we don't do continuous delivery, but all changes that are merged to master and pass CI can get released.
- So if we are releasing daily, how do we build features that are large and naturally take weeks or months to build?
 - Application has notion of "Experiments"
 - Developer creates a feature toggle that changes application behavior
 - toggle is set per-customer
 - Experiments are integrated with our deployment tools
 - We do slow rollout of features to get the right amount of feedback at the right time
 - Experiments are integrated in development workflow
 - Experiments are not an AB testing framework, but sometimes are used as such

qa3

(a40e728bb813cf3f54b805bfb70b3f8adadaed5) (Ops: https://ops.qa.appfolio.com)

**Overview**

Logged in: andrew.mutz

apmbundle**20170313_287883baead5581bddd08d5b4ec5c1d470fd32f**

Committed on: Mon, 13 Mar 2017 11:41:21 vhosts: 47

**20170626_4bbaf924ff56159b2538a14b65e80be912d0a053**

Committed on: Tue, 27 Jun 2017 16:02:58 vhosts: 10



Welcome, **andrew.mutz!** [Application](#) | [Repository](#) | [Host](#) | [Global](#) | [Log Out](#)

Tasks

Destroying 'apmbundle' deployment
"master_a80fb45176c40cc0dc2e2271f669d73901fe0076"
Execution Time: 80 seconds

Modifying vhost 'yunleitest1' by migrating to
"hexRemoveStage1And2_ef8b775f46ddd8dd7a7a897414dae4ab4865e"
Execution Time: 1803 seconds





Engineering Overview: Agile/Scrum

At Appfolio we work in agile teams

- 2-4 developers
- 1 QA Engineer
- 1 product manager
- 1 UX expert

Teams are feature teams, not component teams

- Teams are tasked with solving a problem
- Teams have high degree of freedom to solve problems as they see fit
- No code ownership at Appfolio, teams can make changes to any part of our application

Teams are organized into colleges

- A college has 2-5 teams

Appfolio Property Manager has 5 colleges and 20 teams

Process is heterogeneous





Scaling lessons from industry

- #1 Don't solve the general case
- #2 Keep it simple
- #3 Never forget the customer

Challenges Ahead

- Technology is moving beyond mobile apps and web browsers
 - Package Management
 - Smart Door Locks
 - Voice Telephony
 - Alexa
 - These have different technical problems
 - Vastly different user interface
 - Uptime guarantees need to be much stronger





Challenges Ahead

- Data and Machine Learning
 - Tools are becoming commodity not
 - Our business produces lots of valuable structured data
 - Examples
 - Screening Tenants
 - Understanding Paper Documents
 - Rent Optimization

Challenges Ahead

- Scaling load is sometimes easier than scaling humans
 - Conways law
 - Service Oriented Architectures, Microservices
 - How do you make things maintainable?
 - Does SOA introduce more problems than it solves?



We are hiring full time and interns!

appfolio.com/jobs



Thank you!

Questions?

