

# EVENT-BASED ARCHITECTURE

SCALABLE INTERNET SERVICES  
UCSB - NOV 2021

SEAN MALONEY  
Amazon AWS  
ex Riot Games, ex Appfolio



@SEAN\_SEANNERY



seamalo@amazon.com

# WHO IS THIS GUY?

- AWS CodeCommit Lead
- Riot Games Big Data
- UCSB Lead TA
- US Dept of Labor / Energy

## FUN FACT:

Was a student in this class many years ago. Intern at Appfolio



# EVENT-BASED ARCHITECTURE

---

**1.** THE PROBLEM WITH MICROSERVICES

**2.** EVENT-BASED ARCHITECTURE

**3.** CASE STUDY: RIOT GAMES

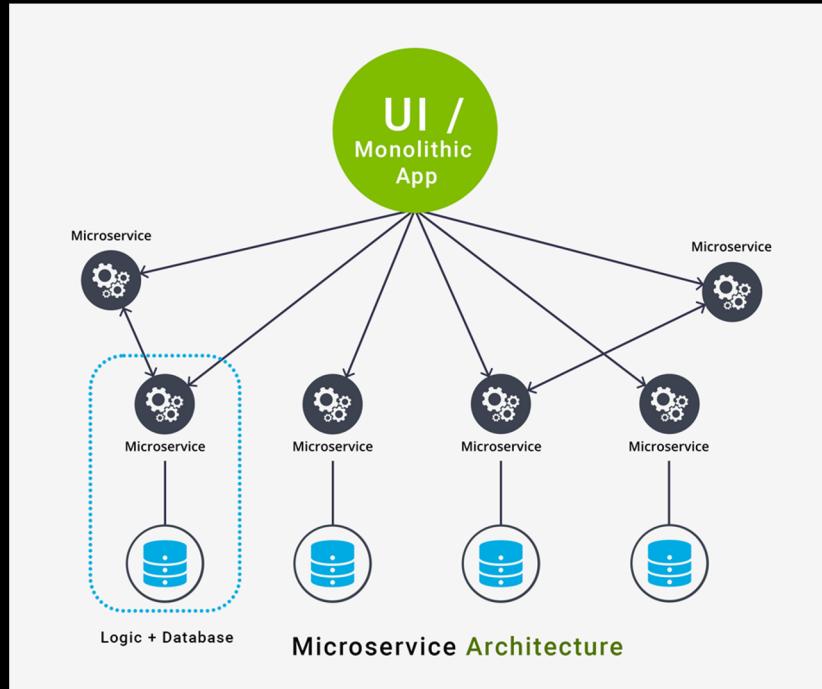
# THE PROBLEM WITH MICROSERVICES



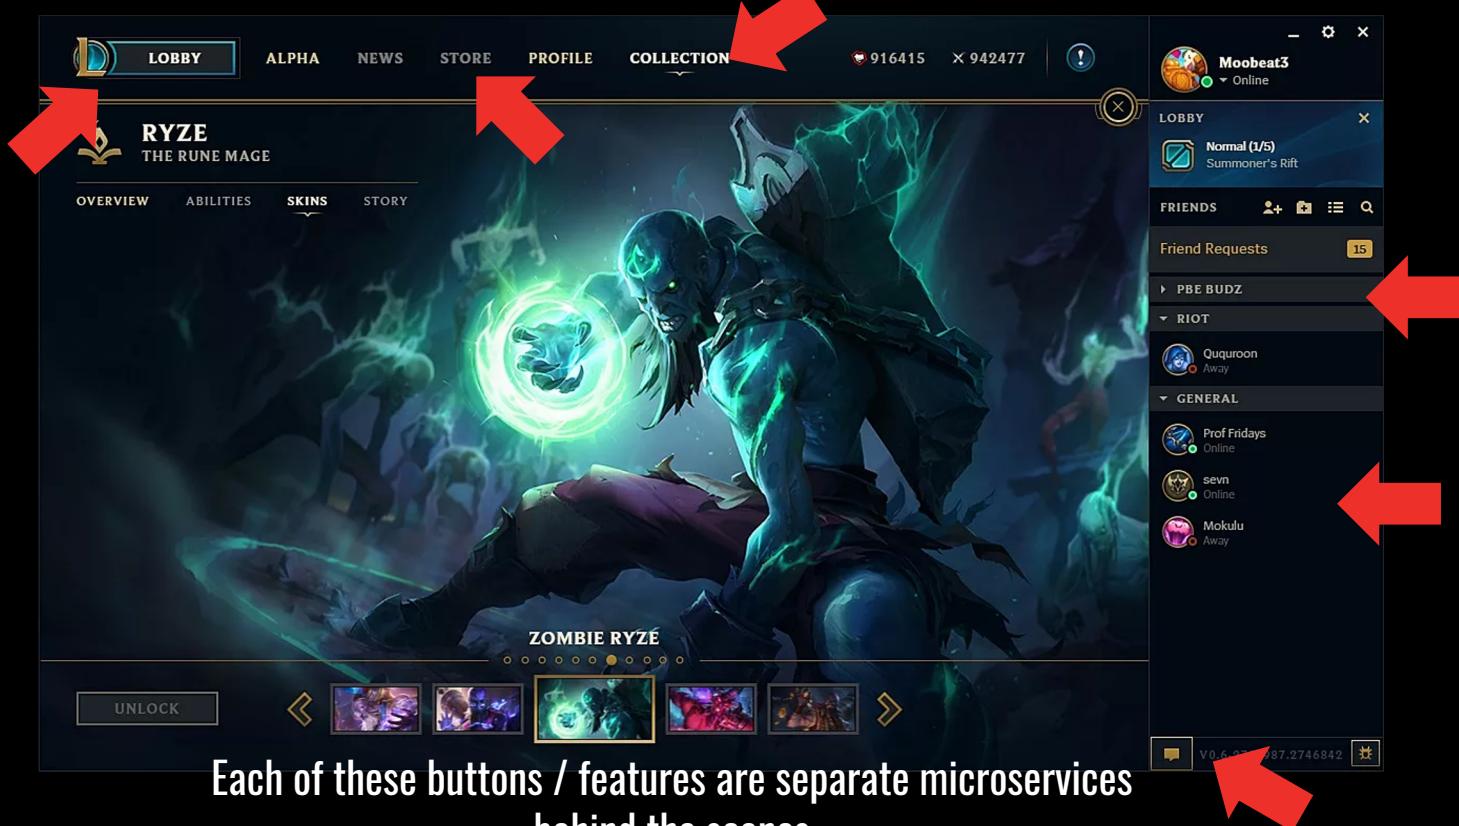
# MICROSERVICE ARCHITECTURE

Loosely Coupled Services responsible for doing one thing well.

Scales well for large enterprises. One team can own one service. Can scale separately for different traffic



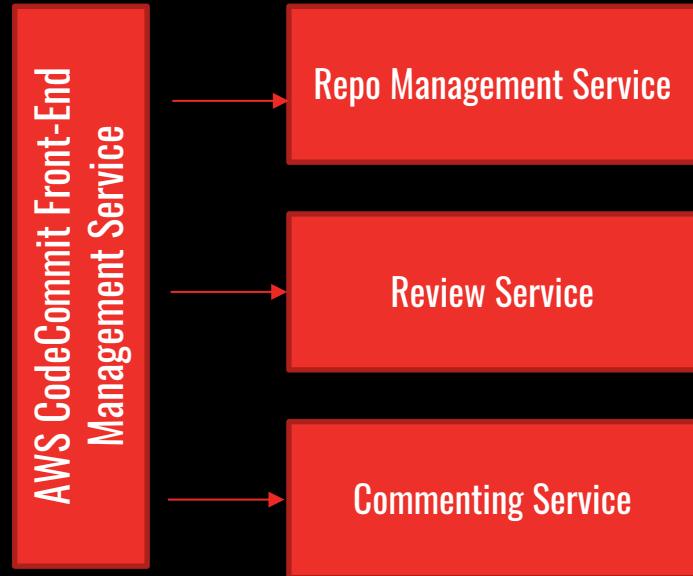
# e.x. League Client



# e.x. AWS CLI

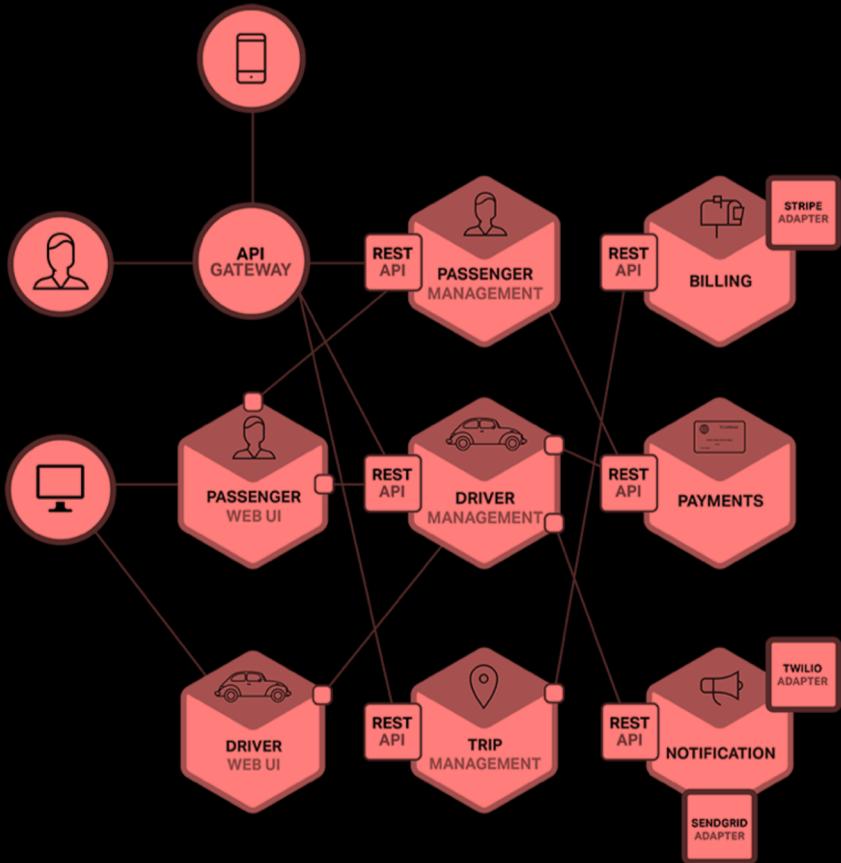
## AWS CodeCommit Command Line Interface:

```
$> aws codecommit [subcmd]  
--create-repository  
--delete-repository  
--get-file  
  
--create-pull-request  
--delete-pull-request  
  
--create-comment  
--delete-comment
```



CodeCommit commands are often handled by separate services.  
In case there is an outage, customers could still do pull requests or comments

# How it Breaks



In a large enterprise, many services have dependencies on many other services

What happens when a service dies, or experiences latency? (like payments svc?)

What happens when a service changes its API structure?

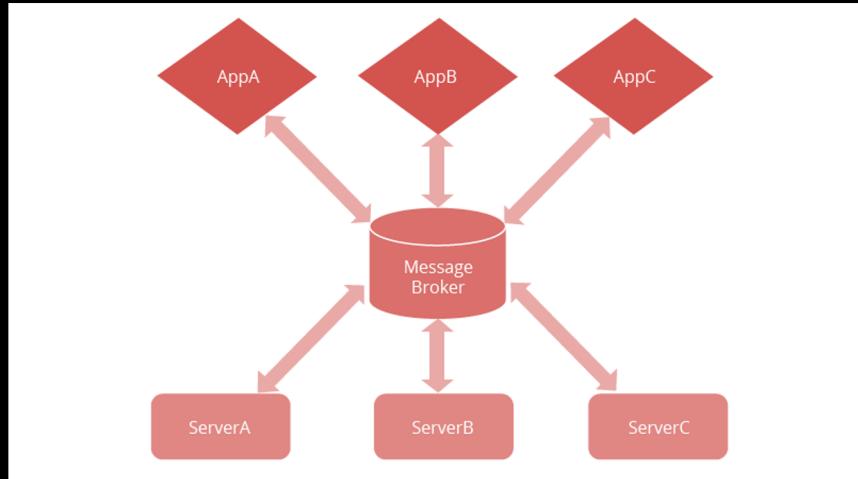
How do I get a unified view of business health?

# EVENT-BASED ARCHITECTURE



# Event-based Architecture

A architectural pattern promoting the production, consumption and processing of events



# Event-based Architecture

HELPS WITH 3 THINGS:

- 1.) Buffering load between services
- 2.) Interoperability / features between services
- 3.) Up-to-date analytics data  
(compared to batch ETLs)

# Event-based Architecture

## What is an event?

Data describing an instance of something happening at a specific point in time

```
"order_event",
{
  "schema" : "order_event_schema_v3",
  "name" : "Bryce Boe",
  "message" : "pumpkin spice latte",
  "store" : "santa_barbara",
  "credit_info" : "a%gGk^d:0ssHjNgs",
  "timestamp" : "2018-10-20"
}
```

```
"employee_timeoff_submission",
{
  "schema" : "emp_pto_schema_v2",
  "name" : "Sean M",
  "startdate" : "2020-12-30",
  "enddate" : "2021-01-02",
  "reason" : "new years vacay",
  "timestamp" : "2018-10-20",
}
```

# Event-based Terminology

**Streams** - an unbounded set of similar events

**Producers** - generate events and send them off

**Consumers** - receive events for consumption

**Processors** - type of consumer that routes or transforms the events and pushes them back into the queue

# A different approach

For example:

Instead of ...

`creditcardservice.purchaseItem()`

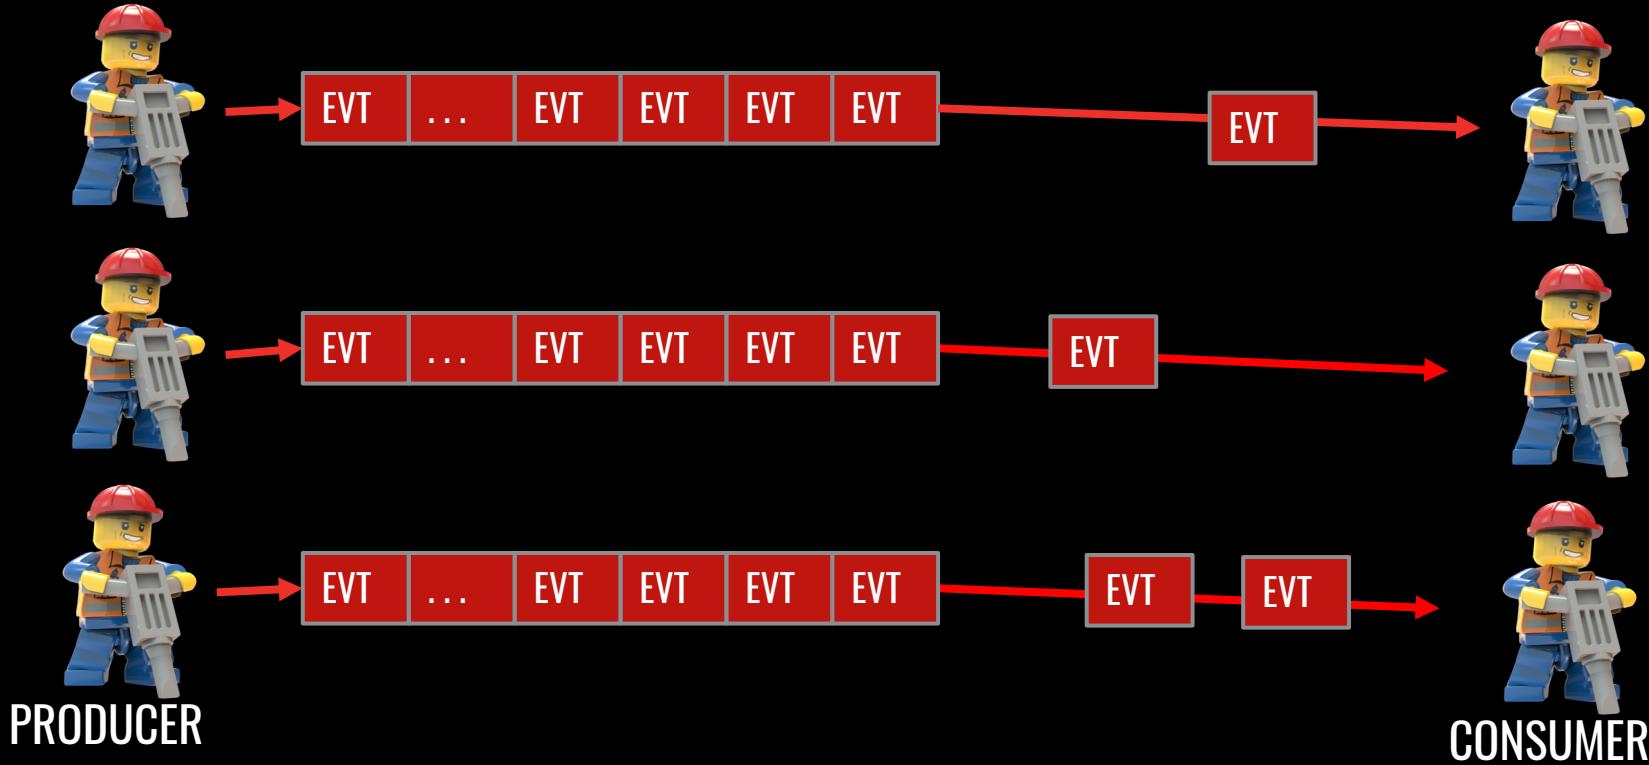
`warehouseservice.updateInventory()`

Do ...

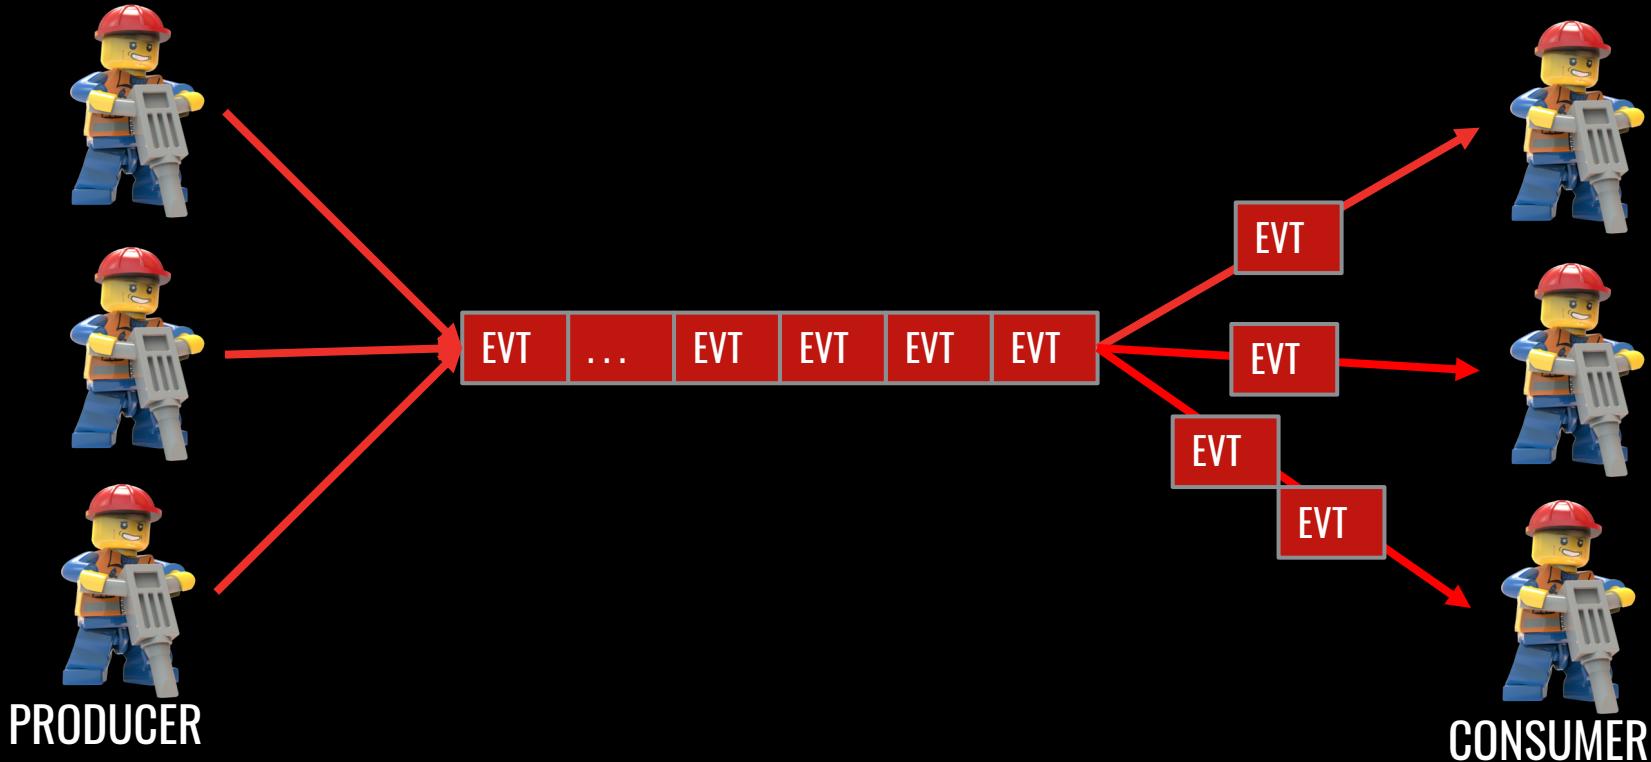
`creditcardservice.processEvent(purchase_event)`

`warehouseservice.processEvent(purchase_event)`

# Message Queues



# Event Bus



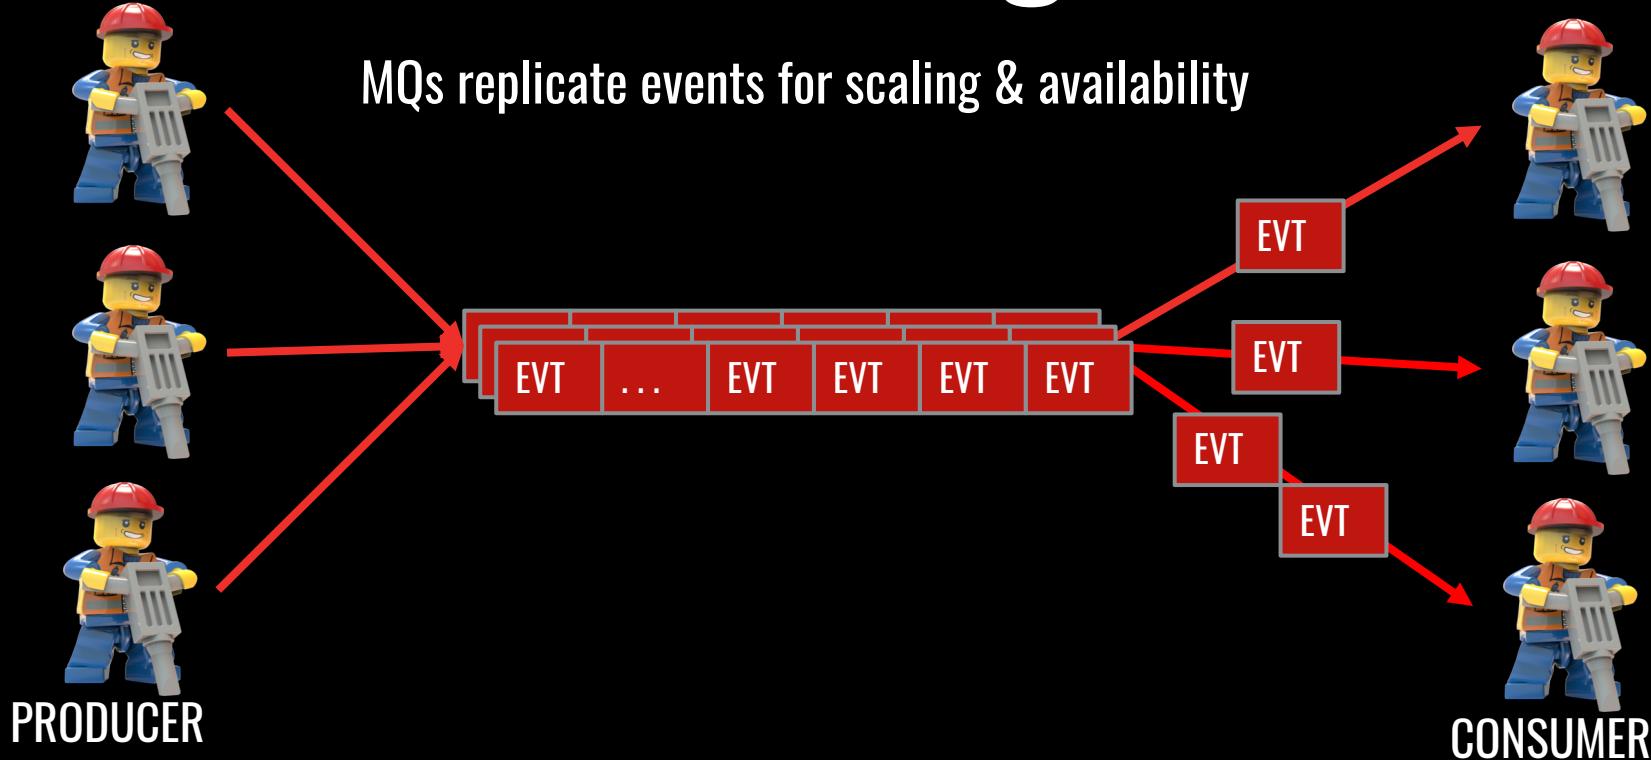
# Message Queues

- Google PubSub
- Amazon Simple Queue Service (SQS)
- Kafka
- RabbitMQ
- Amazon Kenisis
- Microsoft MQ (MSMQ)

# Message Queues

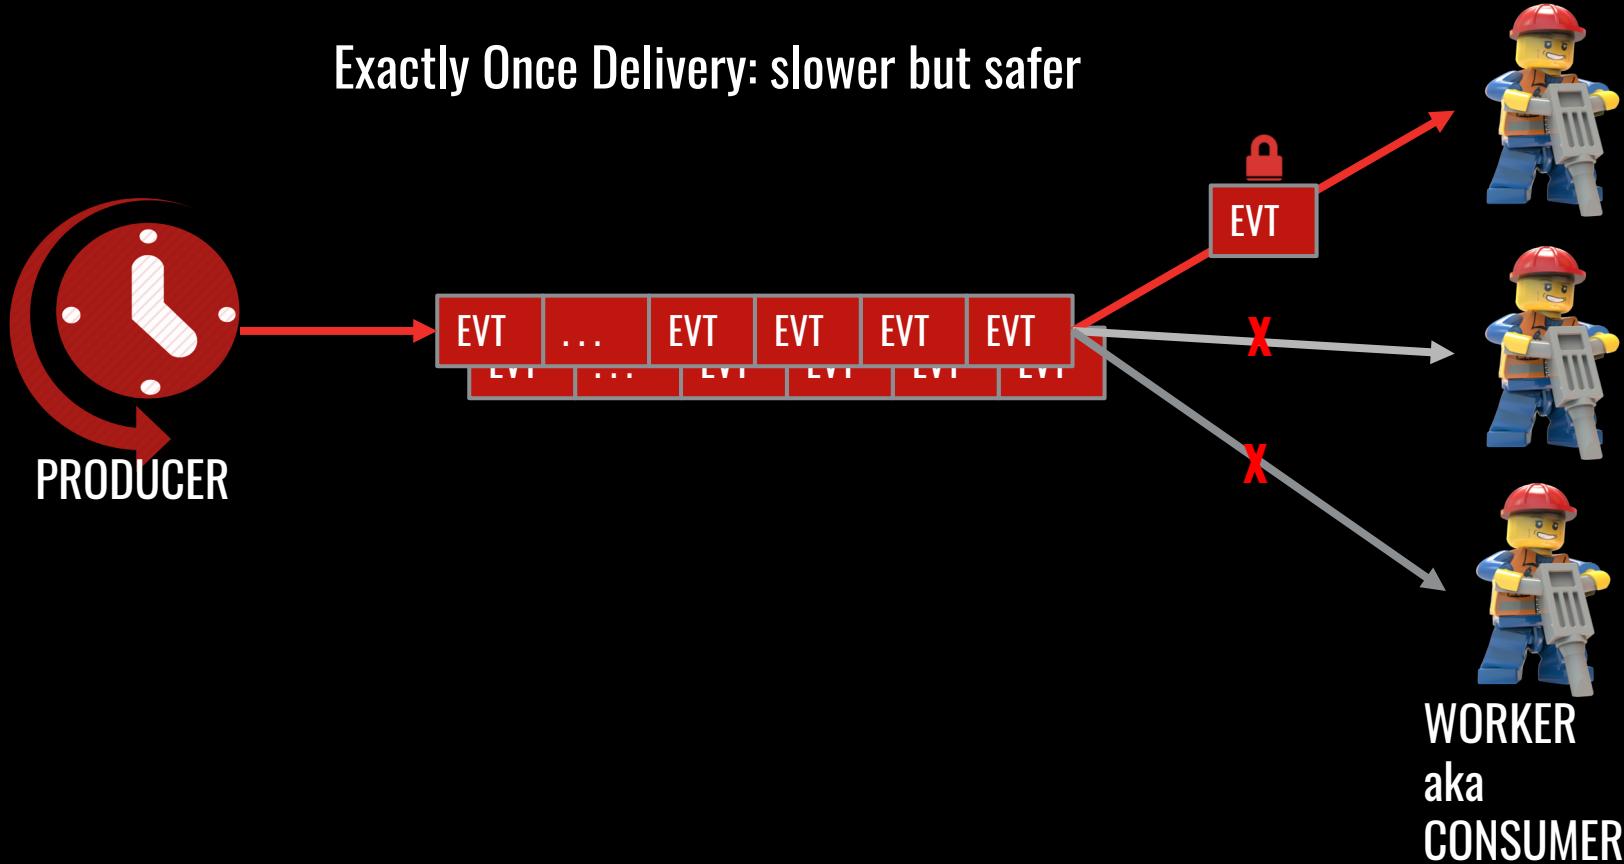
- Redundancy
- Delivery Guarantees
- Easy to Scale
- Asynchronous Communication
- Abstraction / Decoupling

# Scaling



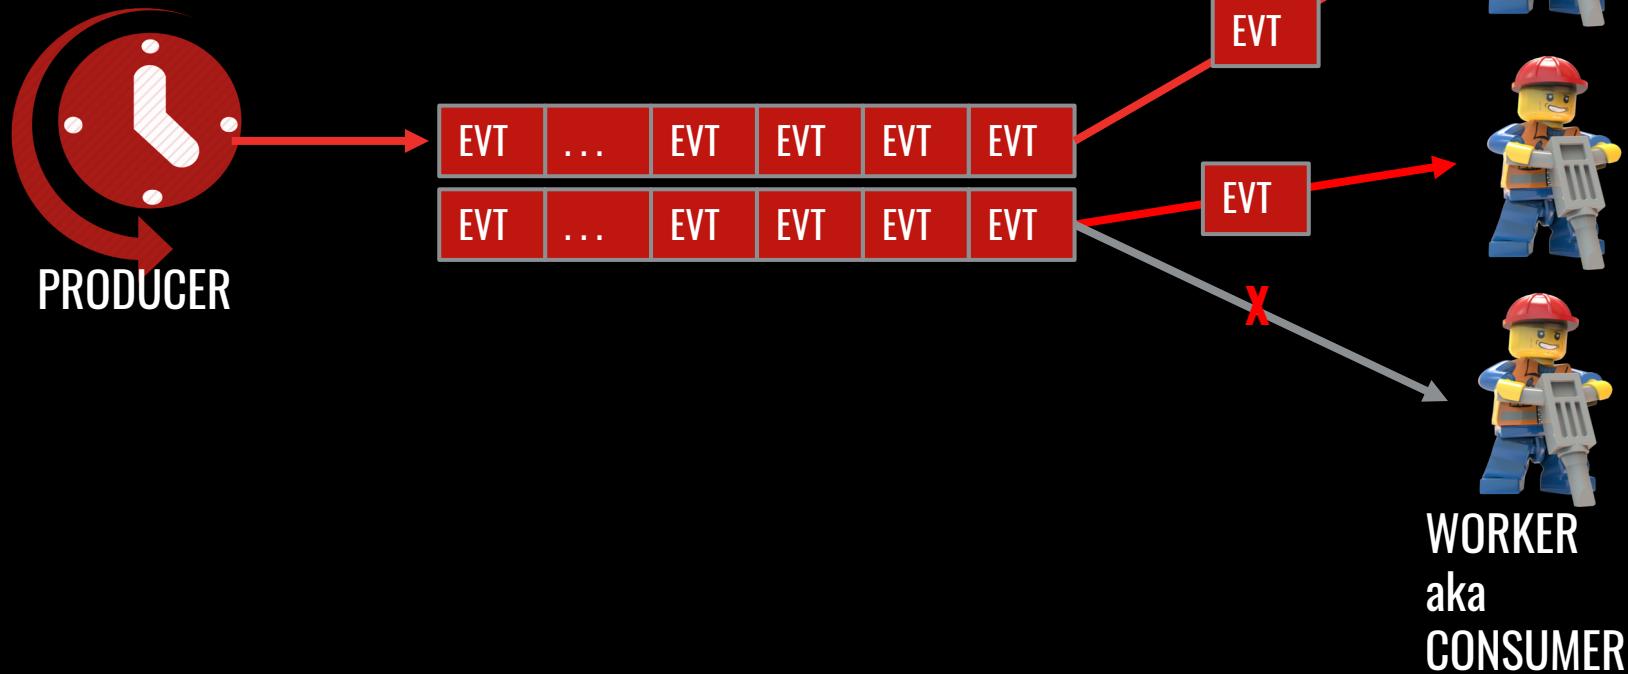
# Delivery Guarantees

Exactly Once Delivery: slower but safer



# Delivery Guarantees

At-least Once Delivery: faster but chance for duplication.



# Idempotency

**Idempotent** – A way to handle duplicate events. An idempotent operation will produce the same results if executed once or multiple times

## EXAMPLE:

Non-Idempotent:

- $x = x * 5;$
- Submitting a purchase

Idempotent:

- $\text{abs}(\text{abs}(x)) = \text{abs}(x)$
- Cancelling a purchase

# Idempotent?

In the transactional OLTP world....

```
INSERT INTO games_played
(SELECT * FROM games_played_na
WHERE date >= '2015-10-25')
```

Potentially with ACID – could get an id already exists

# Idempotent?

In the big data / OLAP world....

```
INSERT INTO games_played
(SELECT * FROM games_played_na
WHERE date >= '2015-10-25')
```

Probably not with noSQL – could get duplicates

# Idempotency

Add application logic to make **idempotent**

```
msg = queue.pop;  
if (processed_games.contains( msg.game_id )  
{  
    return; //do nothing  
else {  
    process_game(msg);  
}
```

# CASE STUDY: RIOT GAMES



# WHAT IS LEAGUE OF LEGENDS?



2009  
LAUNCH

ONLINE  
MULTIPLAYER

WINDOWS  
/ OSX

30-40 MIN  
GAMES

# YOUR CHAMP



THE  
TEAM



THE  
BATTLE  
GROUND



**12 BILLION**  
GAME RELATED EVENTS

**0.5 TRILLION**  
DATA POINTS

**50 TB**  
STORAGE

**DAILY**

**26 PETABYTES**  
PLAYER DATA

**SINCE BETA**

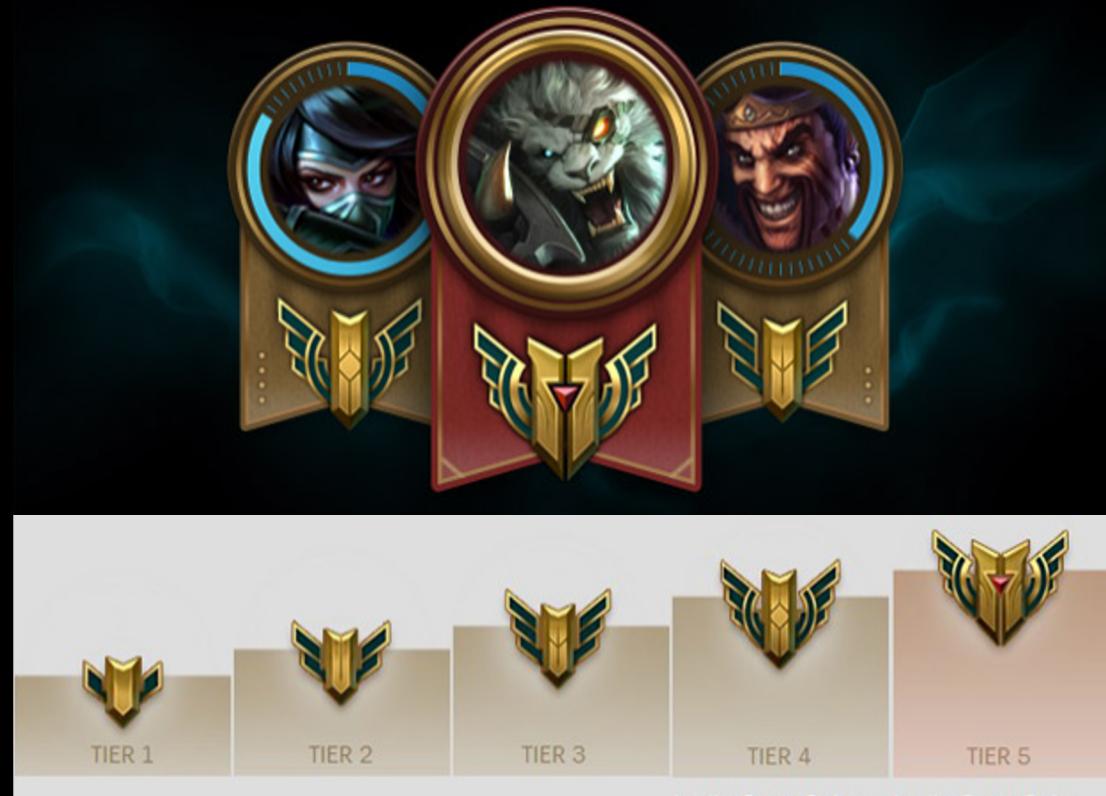
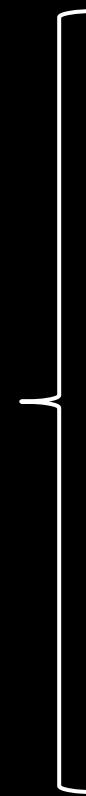
# OFFENSIVE CHAT DETECTION

{ Data science team queries all chat messages in game  
Sentiment analysis and classification

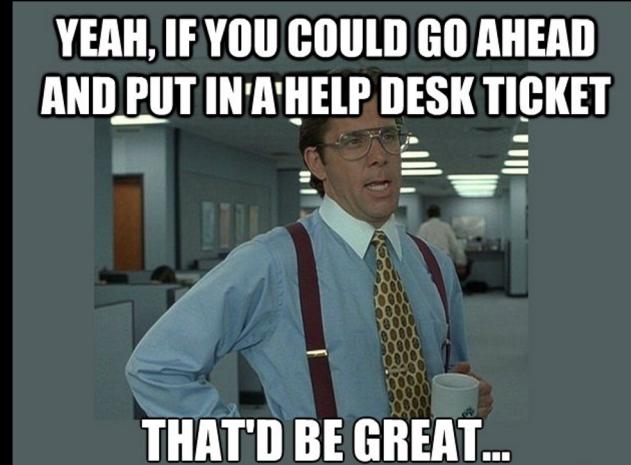
Jônas has ended ezes2396's killing spree! (Bounty: 500G)  
Jônas (Master Yi): sry  
ShadowMaster3000 (Vi): \*\*\*\* your mother \*\*\*\* yi you \*\*\* you noob  
ShadowMaster3000 (Vi): i die and i make ulti and fier and YOPU KILL  
Jônas (Master Yi): :DDDDDD  
ShadowMaster3000 (Vi): i report you \*\*\*\*



# CHAMPION MASTERY



# PLAYER SUPPORT



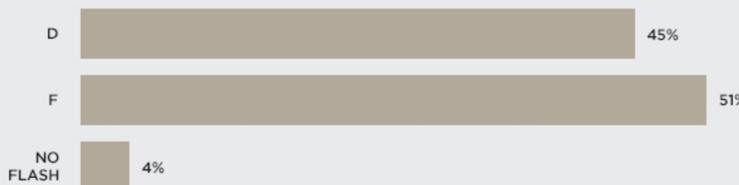
# GAME BALANCE

## FIRST BLOOD RATE BY CHAMPION

RANK	CHAMPION	FIRST BLOOD RATE
1	TALON	24.9%
2	PANTHEON	20.2%
3	KATARINA	19.8%
4	EVELYNN	19.5%
5	LEBLANC	18.6%
6	LEE SIN	17.9%
7	TRYNDAMERE	17.6%

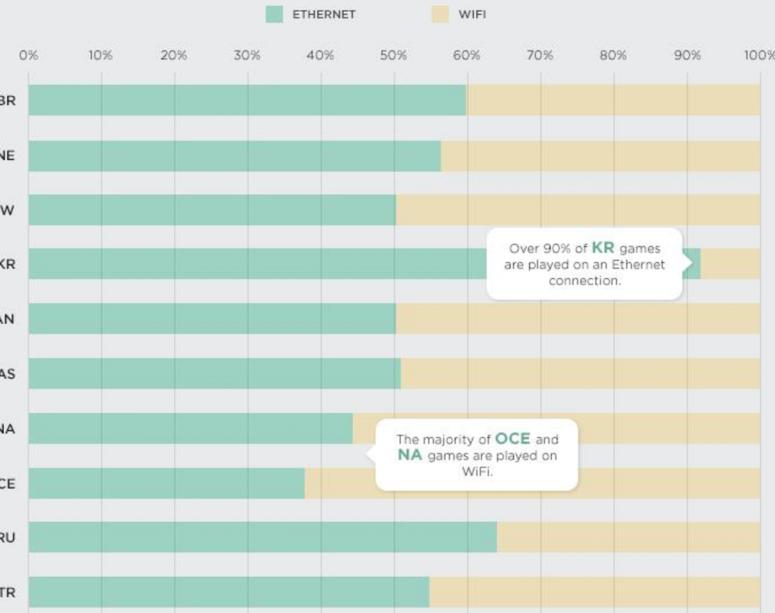
## POPULARITY OF KEYBINDINGS FOR FLASH, WORLDWIDE

BY PERCENTAGE OF RECORDED GAMES

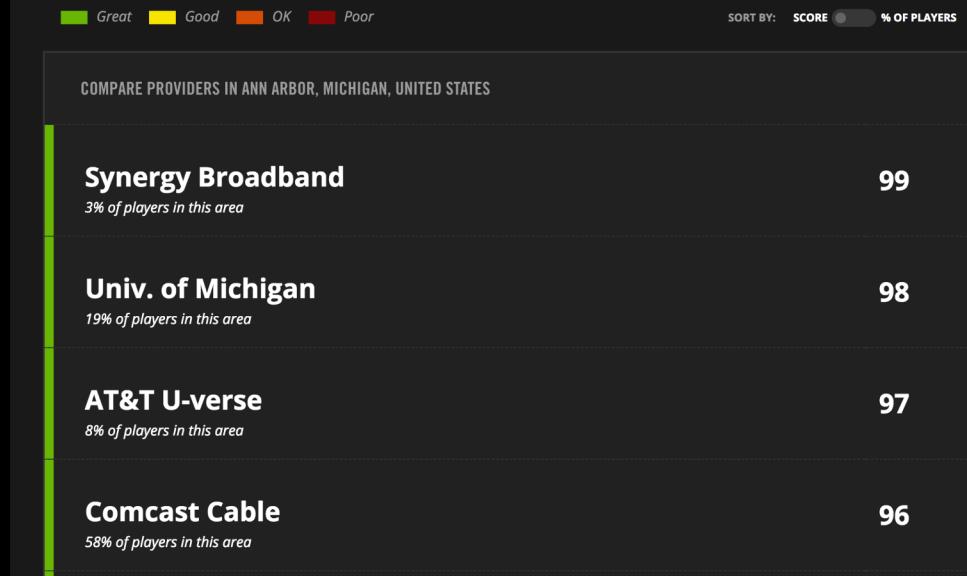


# LATENCY AND NETWORK

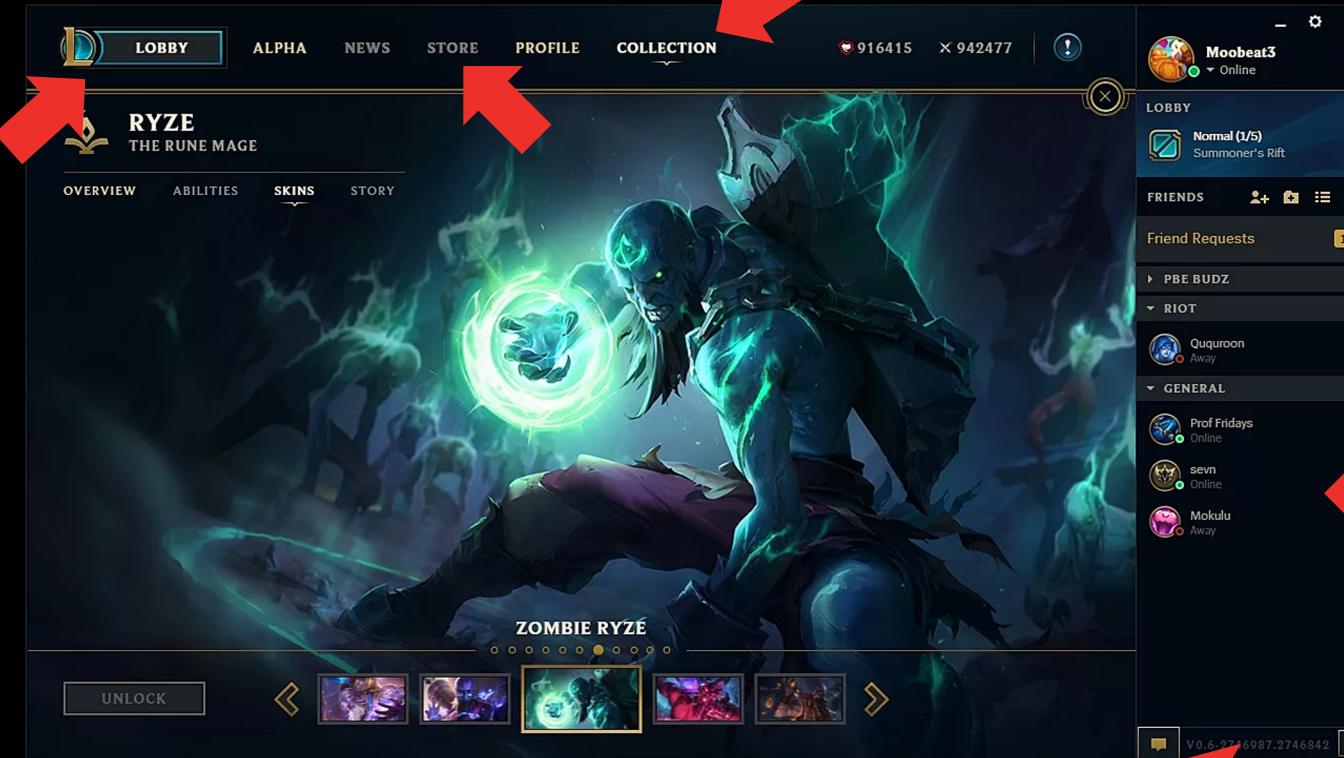
PERCENT OF GAMES PLAYED BY CONNECTION TYPE  
BY REGION



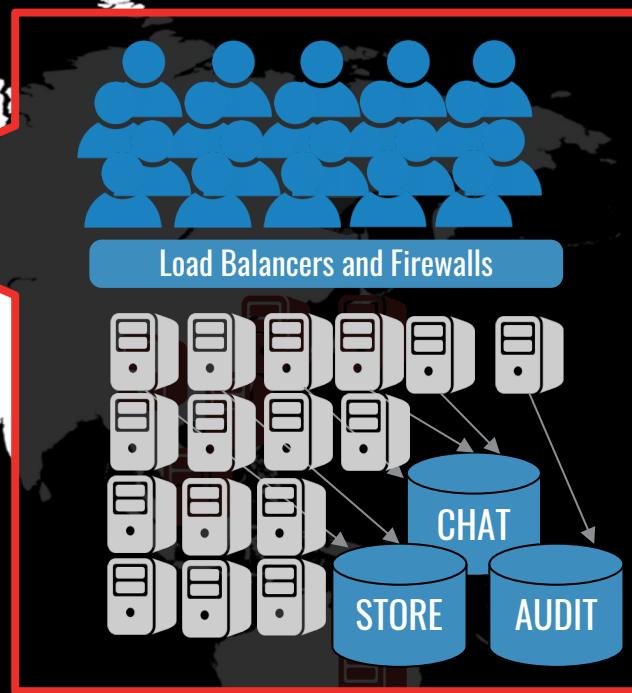
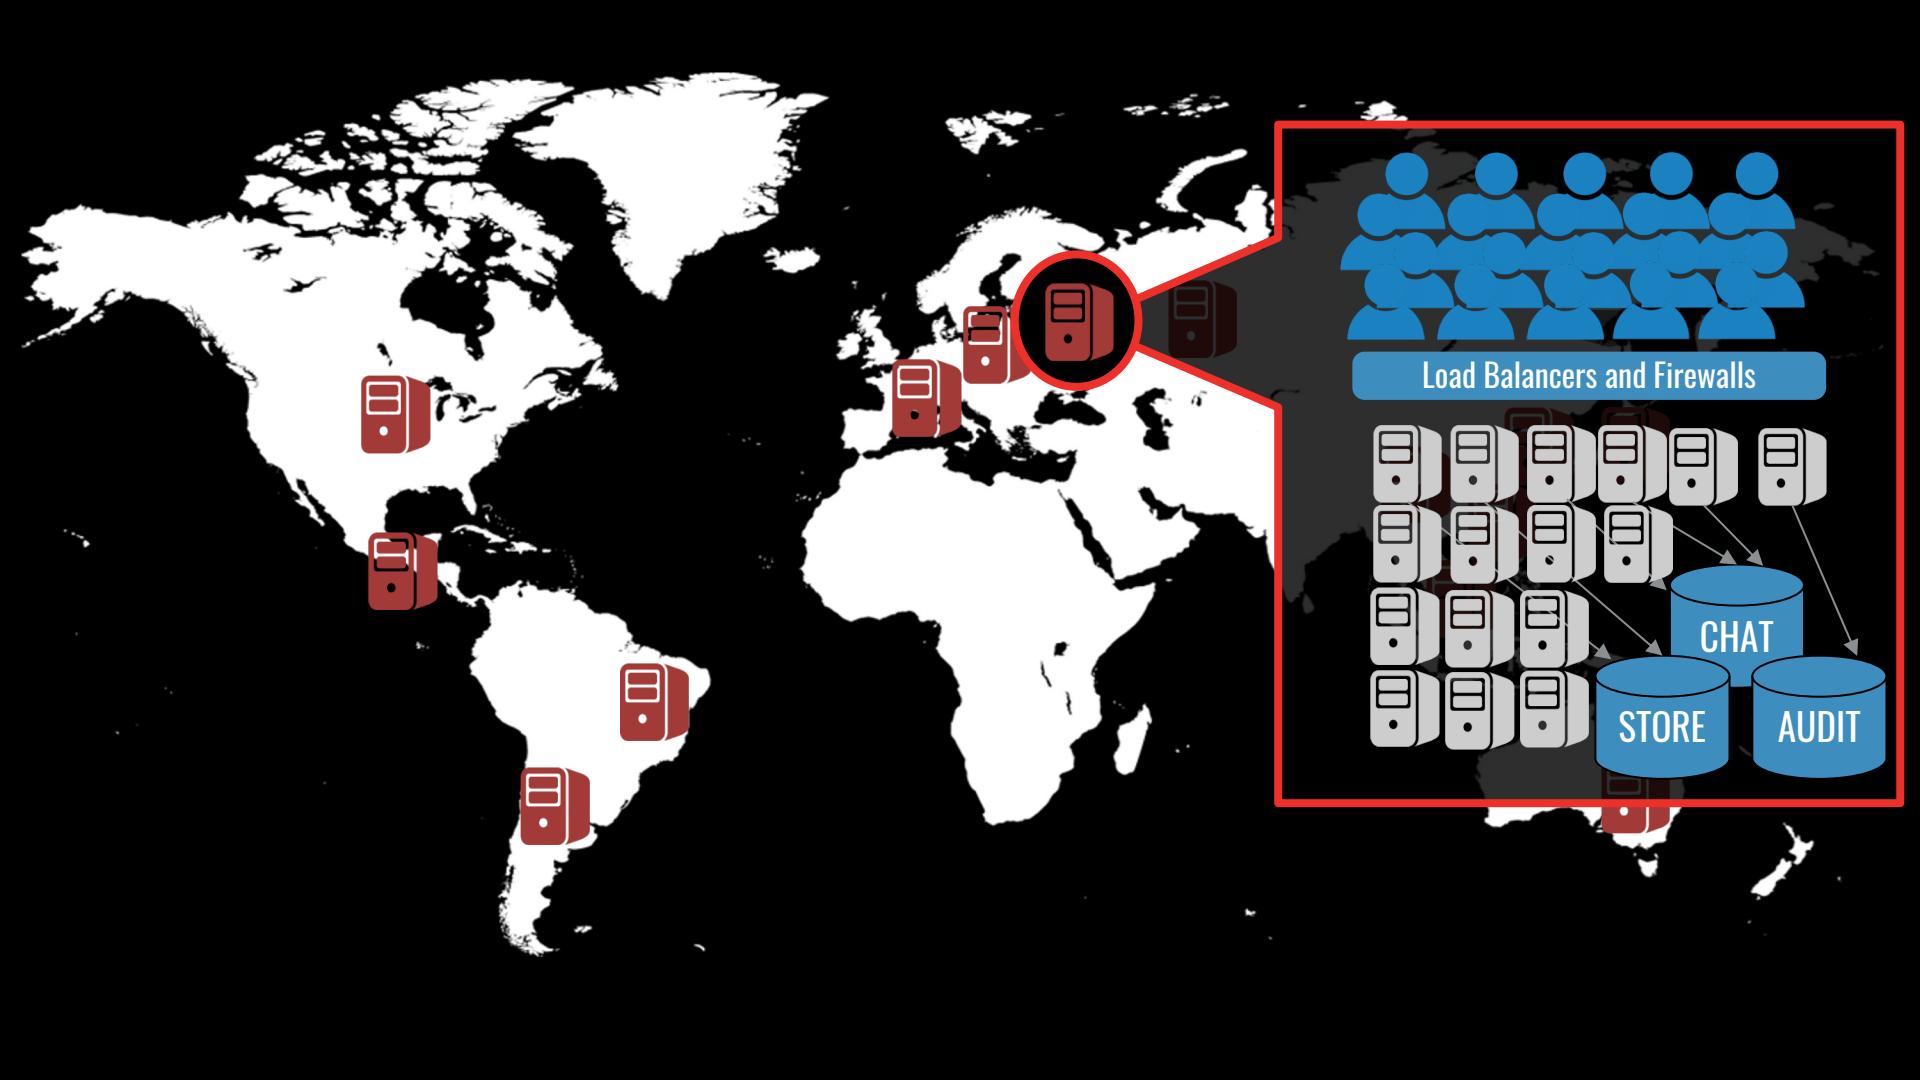
Internet Service Provider Leaderboard [?](#)

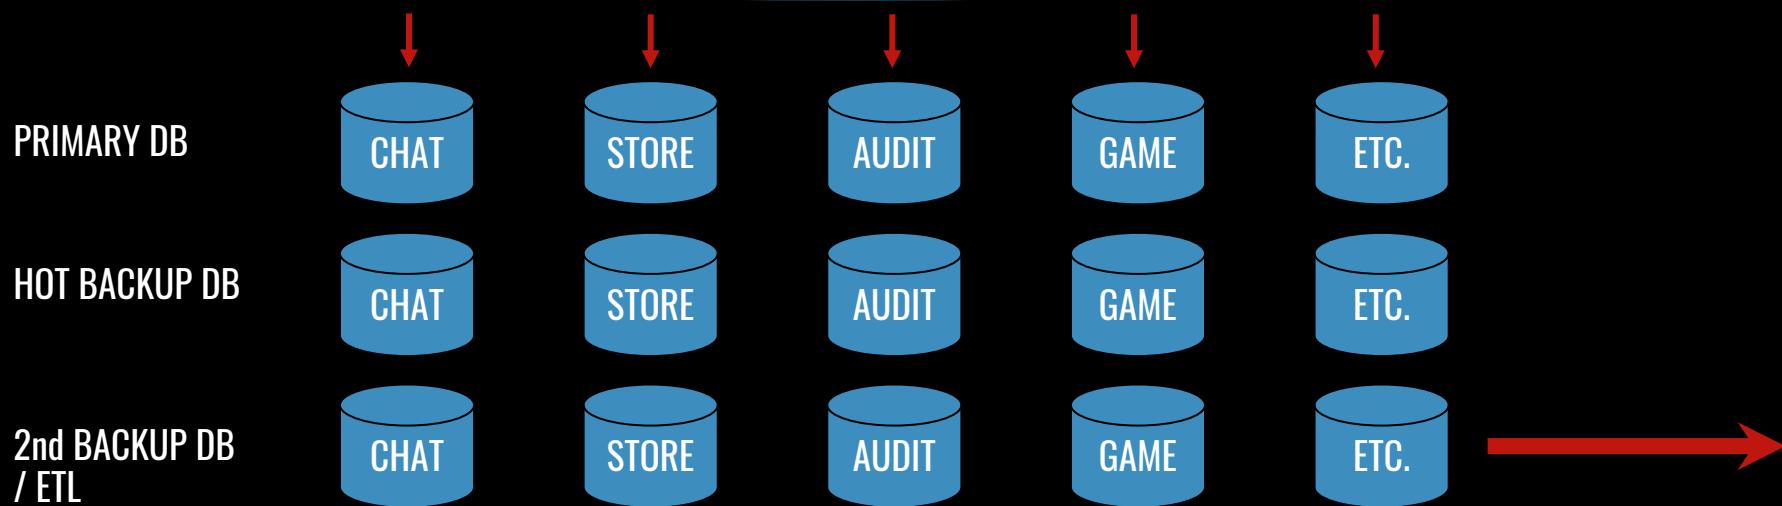
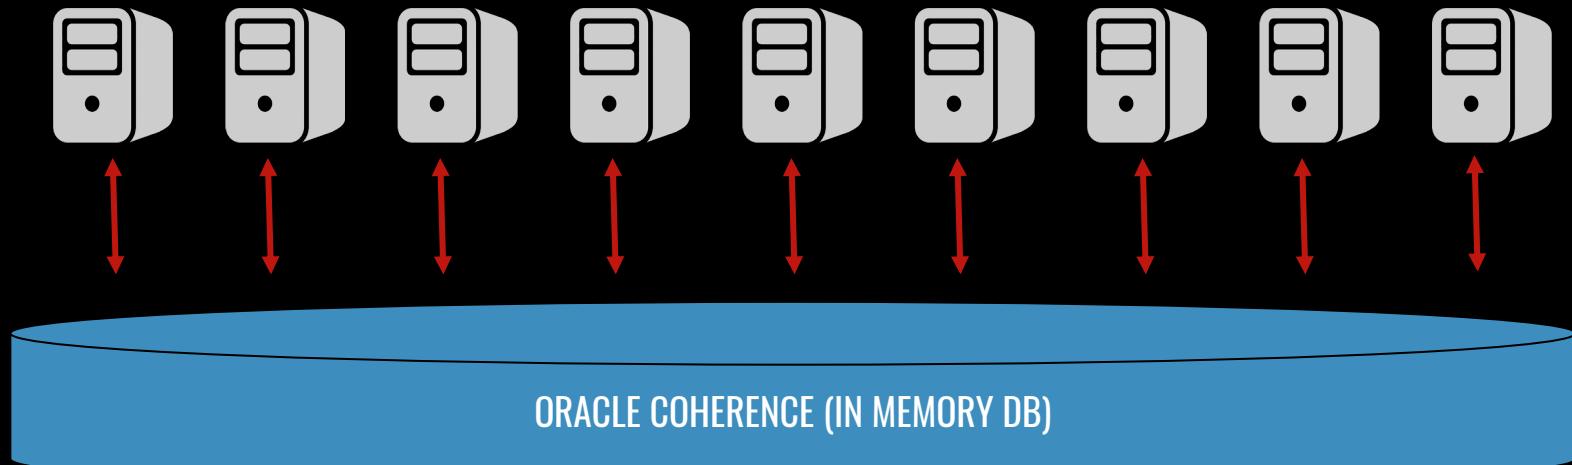


# MICROSERVICES











# EVENTS AT RIOT





Kafka

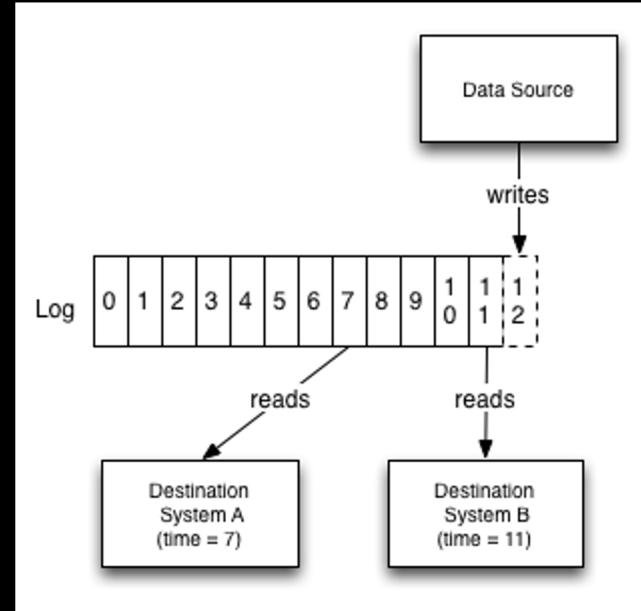
Open-source project maintained by  
Confluent

Very fast distributed event bus

Data is replicated across “partitions” to  
ensure no loss

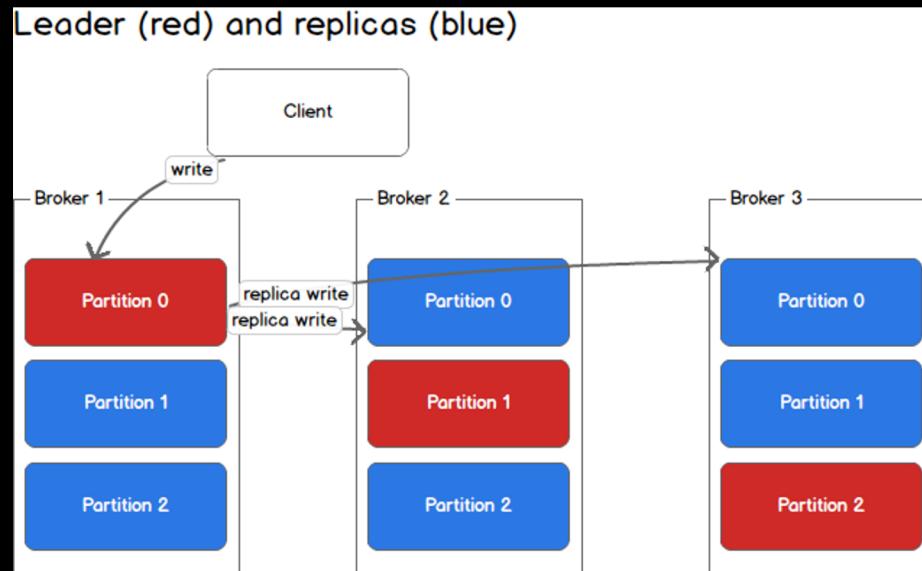
# Kafka

Has a DB Commit Log  
(ooh revolutionary - can replay events)

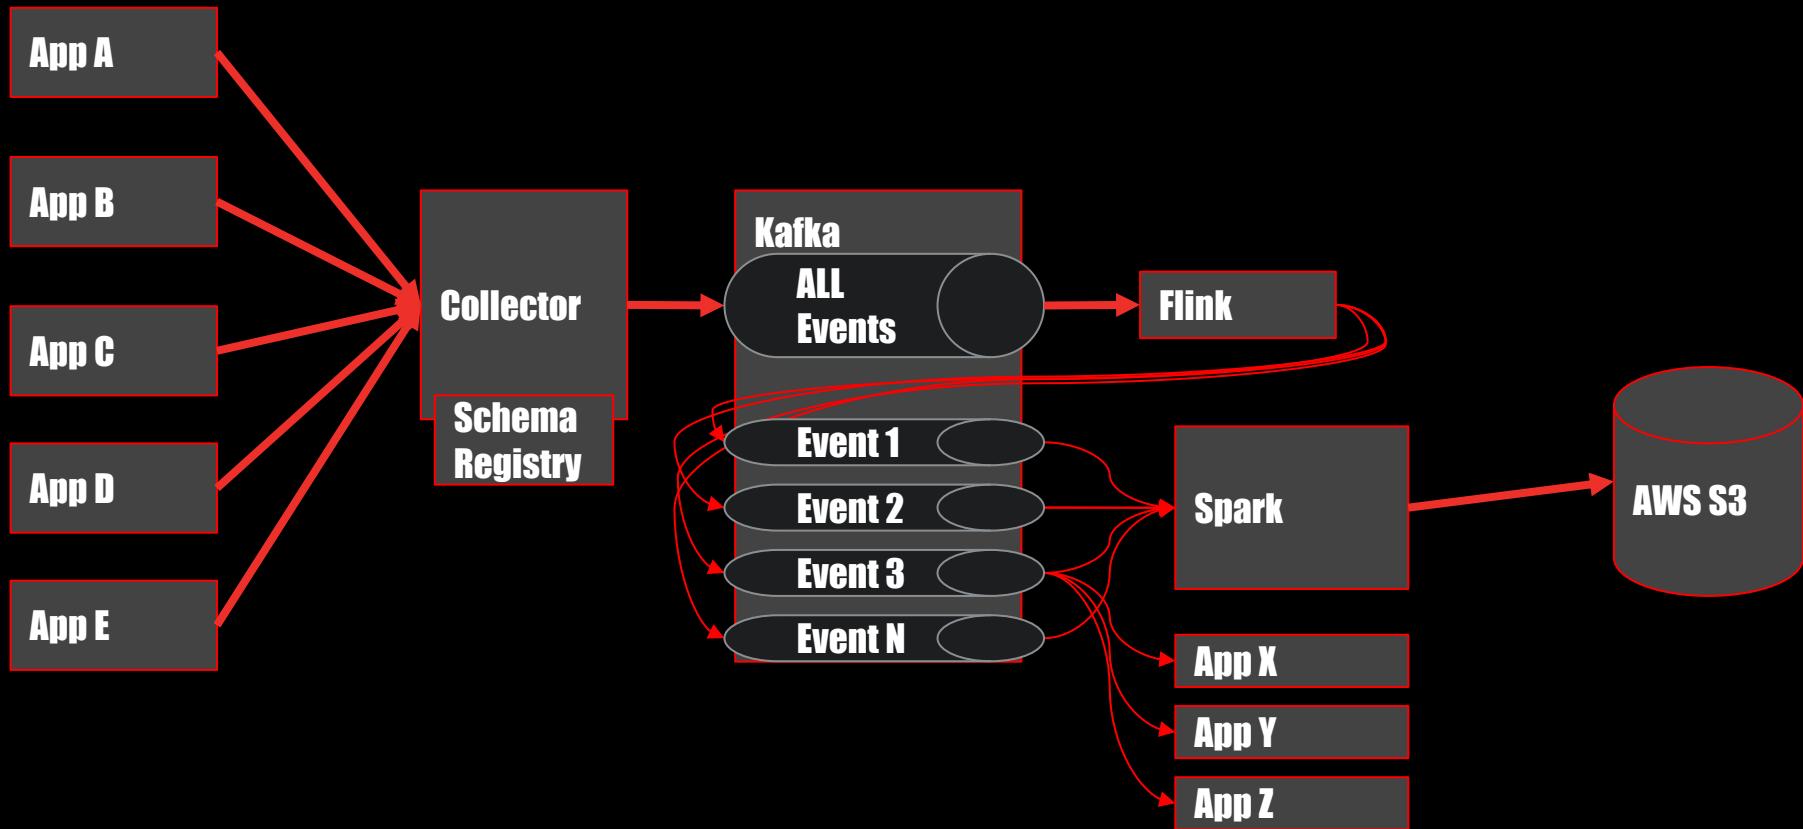


# Kafka

Replicates data and scales for outages



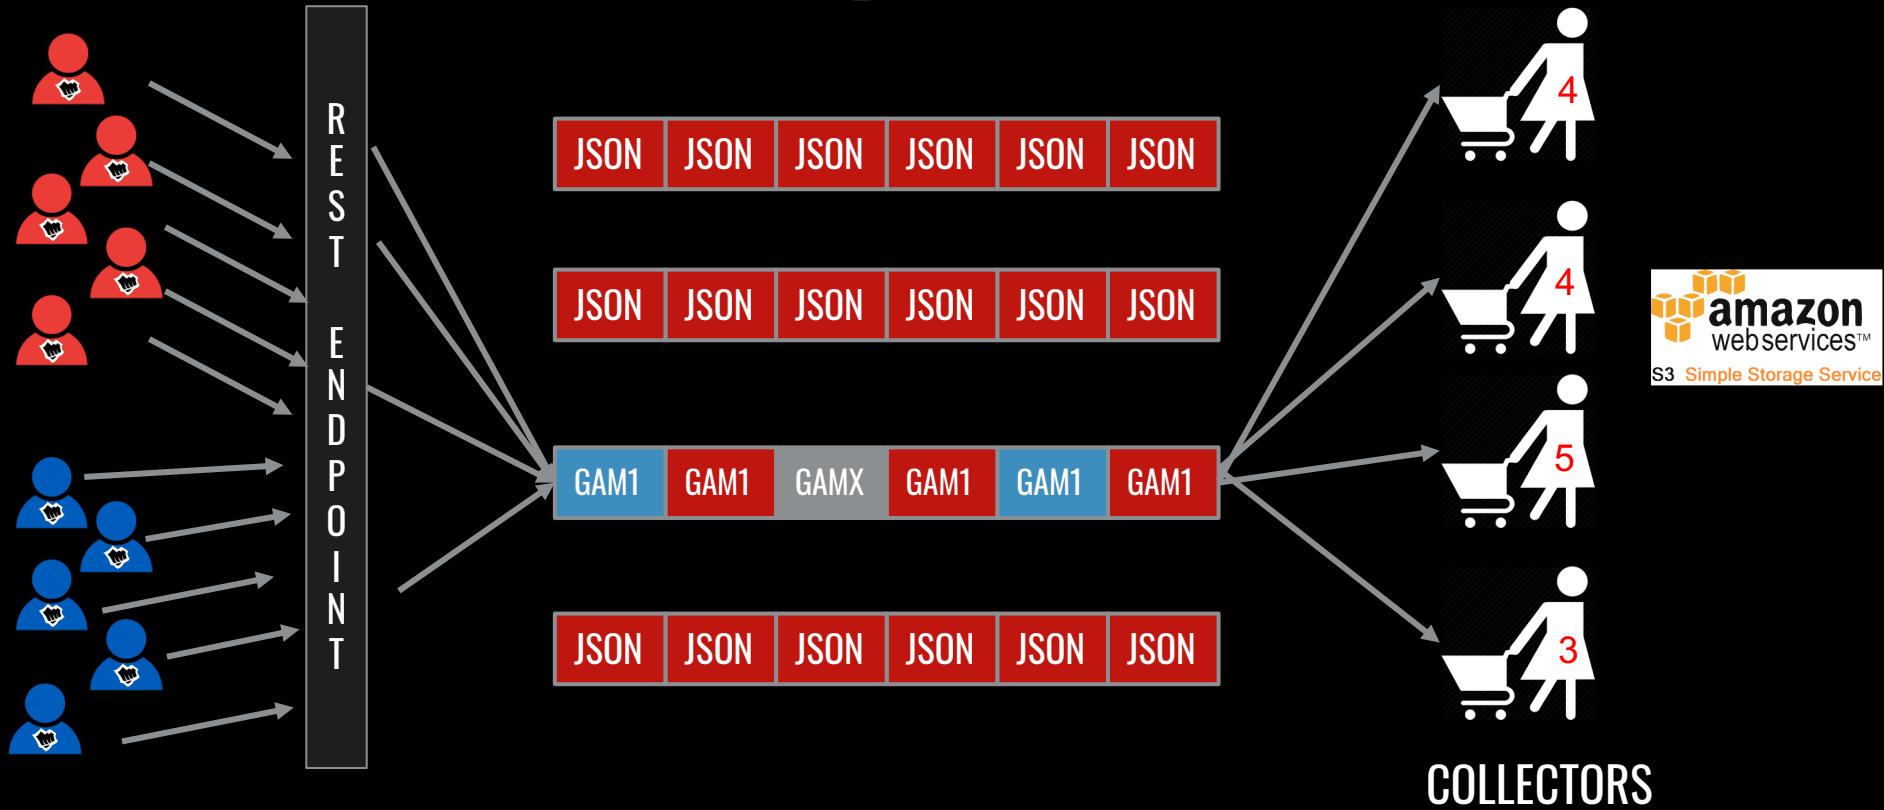
# EVENTS AT RIOT



# Outages

- Network failure between producer and data collection api?  
Producer client a.) exponential backoff retries  
b.) in-memory queue to buffer
- Entire Kafka cluster dies?  
API returns 503 informing client the data wasn't persisted  
and to try again / buffer
- S3 fails, or other consumer?  
Data is buffered for 7 days in Kafka until they recover

# Idempotent?



# QUESTIONS?

SEAN MALONEY

 SEAMALO @  
amazon.com

 @SEAN\_SEANNERY